

Machine Learning

Chapter 10

Machine Learning

- What is learning?

Machine Learning

- Arthur Samuel (1959):
"Field of study that gives computers the ability to learn without being explicitly programmed".
- Tom Mitchell (1997):
"A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E ".

Machine Learning

- How to construct programs that automatically improve with experience.

Machine Learning

- How to construct programs that automatically improve with experience.
- Learning problem:
 - Task **T**
 - Performance measure **P**
 - Training experience **E**

Machine Learning

- Chess game:
 - Task T: playing chess games
 - Performance measure P: percent of games won against opponents
 - Training experience E: playing practice games againsts itself

Machine Learning

- Handwriting recognition:
 - Task **T**: recognizing and classifying handwritten words
 - Performance measure **P**: percent of words correctly classified
 - Training experience **E**: handwritten words with given classifications

Example

Experience

Example	Sky	AirTemp	Humidity	Wind	Water	Forecast	EnjoySport
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Sunny	Warm	High	Strong	Warm	Same	Yes
3	Rainy	Cold	High	Strong	Warm	Change	No
4	Sunny	Warm	High	Strong	Cool	Change	Yes

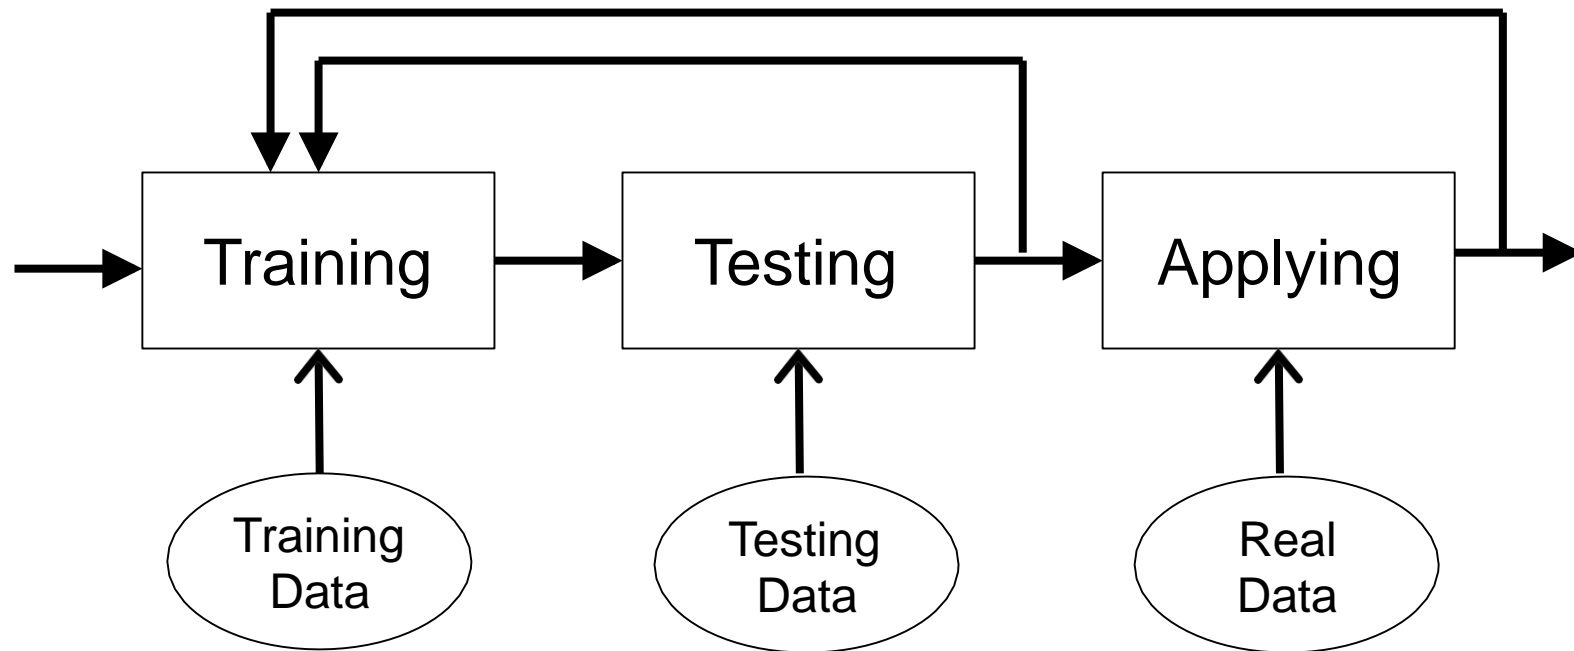
Low

Weak

Prediction

5	Rainy	Cold	High	Strong	Warm	Change	?
6	Sunny	Warm	Normal	Strong	Warm	Same	?
7	Sunny	Warm	Low	Strong	Cool	Same	?

Machine Learning




Machine Learning

- What is learning?

Machine Learning

- What is learning?



Machine Learning

- Learning is an (endless) **generalization** or **induction** process.

Machine Learning

- **Supervised learning**: the learner (learning algorithm) are trained on **labelled** examples, i.e., input where the desired output is known.
- **Unsupervised learning**: the learner operates on **unlabelled** examples, i.e., input where the desired output is unknown.

Concept Learning

- Inferring a boolean-valued function from training examples of its input (**instances**) and output (**classifications**).

Concept Learning

- Learning problem:
 - **Target concept:** a subset of the set of instances X
 $c: X \rightarrow \{0, 1\}$
 - **Target function:**
 $\text{Sky} \times \text{AirTemp} \times \text{Humidity} \times \text{Wind} \times \text{Water} \times \text{Forecast} \rightarrow \{0, 1\}$
 - **Hypothesis:**
Characteristics of all instances of the concept to be learned
 \equiv Constraints on instance attributes
 $h: X \rightarrow \{0, 1\}$

Concept Learning

- Satisfaction:

$h(x) = 1$ iff x satisfies all the constraints of h

$h(x) = 0$ otherwise

- Consistency:

$h(x) = c(x)$ for every instance x of the training examples

- Correctness:

$h(x) = c(x)$ for every instance x of X

Concept Learning

- How to represent a hypothesis?

Concept Learning

- Hypothesis representation (constraints on instance attributes):

<Sky, AirTemp, Humidity, Wind, Water, Forecast>

- ? : any value is acceptable
- single required value
- ∅ : no value is acceptable


- Example:

$h1 = \langle \text{Sunny}, ?, ?, \text{Strong}, ?, ? \rangle$

Concept Learning

- General-to-specific ordering of hypotheses:

$$h_j \geq_g h_k \text{ iff } \forall x \in X: h_k(x) = 1 \Rightarrow h_j(x) = 1$$



Concept Learning

Example	Sky	AirTemp	Humidity	Wind	Water	Forecast	EnjoySport
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Sunny	Warm	High	Strong	Warm	Same	Yes
3	Rainy	Cold	High	Strong	Warm	Change	No
4	Sunny	Warm	High	Strong	Cool	Change	Yes

What is a hypothesis that is consistent with the training examples?

$h = \langle _, _, _, _, _, _ \rangle$

Concept Learning

Example	Sky	AirTemp	Humidity	Wind	Water	Forecast	EnjoySport
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Sunny	Warm	High	Strong	Warm	Same	Yes
3	Rainy	Cold	High	Strong	Warm	Change	No
4	Sunny	Warm	High	Strong	Cool	Change	Yes

What is the **most specific** hypothesis that is consistent with the training examples?

h = < _, _ , _ , _ , _ , _ >

FIND-S

Example	Sky	AirTemp	Humidity	Wind	Water	Forecast	EnjoySport
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Sunny	Warm	High	Strong	Warm	Same	Yes
3	Rainy	Cold	High	Strong	Warm	Change	No
4	Sunny	Warm	High	Strong	Cool	Change	Yes

$h = \langle \emptyset, \emptyset, \emptyset, \emptyset, \emptyset, \emptyset, \emptyset \rangle$

$h = \langle \text{Sunny}, \text{Warm}, \text{Normal}, \text{Strong}, \text{Warm}, \text{Same} \rangle$

$h = \langle \text{Sunny}, \text{Warm}, \quad ? \quad, \text{Strong}, \text{Warm}, \text{Same} \rangle$

$h = \langle \text{Sunny}, \text{Warm}, \quad ? \quad, \text{Strong}, \quad ? \quad, \quad ? \quad \rangle$

FIND-S

- Initialize h to the most specific hypothesis in H :
- For each positive training instance x :
 - For each attribute constraint a_i in h :
 - If the constraint is not satisfied by x
 - Then replace a_i by the next more general constraint satisfied by x
- Output hypothesis h

FIND-S

Example	Sky	AirTemp	Humidity	Wind	Water	Forecast	EnjoySport
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Sunny	Warm	High	Strong	Warm	Same	Yes
3	Rainy	Cold	High	Strong	Warm	Change	No
4	Sunny	Warm	High	Strong	Cool	Change	Yes

$h = \langle \text{Sunny, Warm, } ? , \text{Strong, } ? , ? \rangle$

Prediction

5	Rainy	Cold	High	Strong	Warm	Change	No
6	Sunny	Warm	Normal	Strong	Warm	Same	Yes
7	Sunny	Warm	Low	Strong	Cool	Same	Yes

FIND-S

- The output hypothesis is the **most specific** one that satisfies all positive training examples.

FIND-S

- The result is consistent with the **positive** training examples.

FIND-S

- Is the result is consistent with the **negative** training examples?

FIND-S

Example	Sky	AirTemp	Humidity	Wind	Water	Forecast	EnjoySport
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Sunny	Warm	High	Strong	Warm	Same	Yes
3	Rainy	Cold	High	Strong	Warm	Change	No
4	Sunny	Warm	High	Strong	Cool	Change	Yes
5	Sunny	Warm	Normal	Strong	Cool	Change	No

h = <Sunny, Warm, ? , Strong, ? , ? >

FIND-S

- The result is consistent with the **negative** training examples if the **target concept** is contained in **H** (and the training examples are correct).

FIND-S

- The result is consistent with the **negative** training examples if the **target concept** is contained in **H** (and the training examples are correct).
 - Sizes of the space:
 - Size of the instance space: $|X| = 3 \cdot 2 \cdot 2 \cdot 2 \cdot 2 \cdot 2 = 96$
 - Size of the concept space $C = 2^{|X|} = 2^{96}$
 - Size of the hypothesis space $H = (4 \cdot 3 \cdot 3 \cdot 3 \cdot 3 \cdot 3) + 1 = 973 \ll 2^{96}$
- ⇒ The target concept (in **C**) may not be contained in **H**.

FIND-S

- Questions:
 - Has the learner **converged** to the target concept, as there can be several consistent hypotheses (with both positive and negative training examples)?
 - Why the **most specific** hypothesis is preferred?
 - What if there are **several maximally specific** consistent hypotheses?
 - What if the training examples are **not correct**?

List-then-Eliminate Algorithm

- **Version space**: a set of all hypotheses that are consistent with the training examples.
- Algorithm:
 - Initial version space = set containing every hypothesis in H
 - For each **training example** $\langle x, c(x) \rangle$, remove from the version space any hypothesis h for which $h(x) \neq c(x)$
 - Output the hypotheses in the version space

List-then-Eliminate Algorithm

- Requires an exhaustive enumeration of all hypotheses in H

Compact Representation of Version Space

- **G** (the generic boundary): set of the most generic hypotheses of **H** consistent with the training data **D**:


$$G = \{g \in H \mid \text{consistent}(g, D) \wedge \neg \exists g' \in H: g' >_g g \wedge \text{consistent}(g', D)\}$$

- **S** (the specific boundary): set of the most specific hypotheses of **H** consistent with the training data **D**:

$$S = \{s \in H \mid \text{consistent}(s, D) \wedge \neg \exists s' \in H: s >_g s' \wedge \text{consistent}(s', D)\}$$

Compact Representation of Version Space

- Version space = $\langle G, S \rangle = \{h \in H \mid \exists g \in G \exists s \in S: g \geq_g h \geq_g s\}$



Candidate-Elimination Algorithm

Example	Sky	AirTemp	Humidity	Wind	Water	Forecast	EnjoySport
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Sunny	Warm	High	Strong	Warm	Same	Yes
3	Rainy	Cold	High	Strong	Warm	Change	No
4	Sunny	Warm	High	Strong	Cool	Change	Yes

$S_0 = \{ \langle \emptyset, \emptyset, \emptyset, \emptyset, \emptyset, \emptyset \rangle \}$

$G_0 = \{ \langle ?, ?, ?, ?, ?, ? \rangle \}$

$S_1 = \{ \langle \text{Sunny}, \text{Warm}, \text{Normal}, \text{Strong}, \text{Warm}, \text{Same} \rangle \}$

$G_1 = \{ \langle ?, ?, ?, ?, ?, ? \rangle \}$

$S_2 = \{ \langle \text{Sunny}, \text{Warm}, ?, \text{Strong}, \text{Warm}, \text{Same} \rangle \}$


$G_2 = \{ \langle ?, ?, ?, ?, ?, ? \rangle \}$

$S_3 = \{ \langle \text{Sunny}, \text{Warm}, ?, \text{Strong}, \text{Warm}, \text{Same} \rangle \}$


$G_3 = \{ \langle \text{Sunny}, ?, ?, ?, ?, ? \rangle, \langle ?, \text{Warm}, ?, ?, ?, ? \rangle, \langle ?, ?, ?, ?, ?, \text{Same} \rangle \}$

$S_4 = \{ \langle \text{Sunny}, \text{Warm}, ?, \text{Strong}, ?, ? \rangle \}$

$G_4 = \{ \langle \text{Sunny}, ?, ?, ?, ?, ? \rangle, \langle ?, \text{Warm}, ?, ?, ?, ? \rangle \}$



Candidate-Elimination Algorithm



Candidate-Elimination Algorithm

- Initialize G to the set of **maximally general** hypotheses in H
- Initialize S to the set of **maximally specific** hypotheses in H

Candidate-Elimination Algorithm

- For each positive example d :
 - Remove from G any hypothesis inconsistent with d
 - For each s in S that is inconsistent with d :
 - Remove s from S
 - Add to S all least generalizations h of s , such that h is consistent with d and some hypothesis in G is more general than h
 - Remove from S any hypothesis that is more general than another hypothesis in S

Candidate-Elimination Algorithm


- For each negative example d :
 - Remove from S any hypothesis inconsistent with d
 - For each g in G that is inconsistent with d :
 - Remove g from G
 - Add to G all least specializations h of g , such that h is consistent with d and some hypothesis in S is more specific than h
 - Remove from G any hypothesis that is more specific than another hypothesis in G

Candidate-Elimination Algorithm

- The version space will **converge** toward the correct target concepts if:
 - **H** contains the correct target concept
 - There are no errors in the training examples
- A training instance to be **requested next** should discriminate among the alternative hypotheses in the current version space:

Candidate-Elimination Algorithm

- Partially learned concept can be used to classify new instances using the majority rule.



5	Rainy	Warm	High	Strong	Cool	Same	?
---	-------	------	------	--------	------	------	---

Inductive Bias

- Size of the instance space: $|X| = 3 \cdot 2 \cdot 2 \cdot 2 \cdot 2 \cdot 2 = 96$
- Number of possible concepts = $2^{|X|} = 2^{96}$
- Size of $H = (4 \cdot 3 \cdot 3 \cdot 3 \cdot 3 \cdot 3) + 1 = 973 \ll 2^{96}$

Inductive Bias

- Size of the instance space: $|X| = 3 \cdot 2 \cdot 2 \cdot 2 \cdot 2 \cdot 2 = 96$
- Number of possible concepts = $2^{|X|} = 2^{96}$
- Size of $H = (4 \cdot 3 \cdot 3 \cdot 3 \cdot 3 \cdot 3) + 1 = 973 \ll 2^{96}$

\Rightarrow a **biased** hypothesis space


Inductive Bias

- An **unbiased** hypothesis space H' that can represent every subset of the instance space X : **Propositional logic sentences**
- Positive examples: x_1, x_2, x_3
Negative examples: x_4, x_5

$$h(x) \equiv (x = x_1) \vee (x = x_2) \vee (x = x_3) \equiv x_1 \vee x_2 \vee x_3$$

$$h'(x) \equiv (x \neq x_4) \wedge (x \neq x_5) \equiv \neg x_4 \wedge \neg x_5$$

Inductive Bias



Any new instance x is classified positive **by half** of the version space, and negative by the other half

\Rightarrow **not classifiable**

Inductive Bias

Example	Day	Actor	Price	EasyTicket
1	Mon	Famous	Expensive	Yes
2	Sat	Famous	Moderate	No
3	Sun	Infamous	Cheap	No
4	Wed	Infamous	Moderate	Yes
5	Sun	Famous	Expensive	No
6	Thu	Infamous	Cheap	Yes
7	Tue	Famous	Expensive	Yes
8	Sat	Famous	Cheap	No

9	Wed	Famous	Cheap	?
10	Sat	Infamous	Expensive	?

Inductive Bias

Example	Quality	Price	Buy
1	Good	Low	Yes
2	Bad	High	No

3	Good	High	?
4	Bad	Low	?

Inductive Bias

- A learner that makes no prior assumptions regarding the identity of the target concept cannot classify any unseen instances.

Homework

Exercises

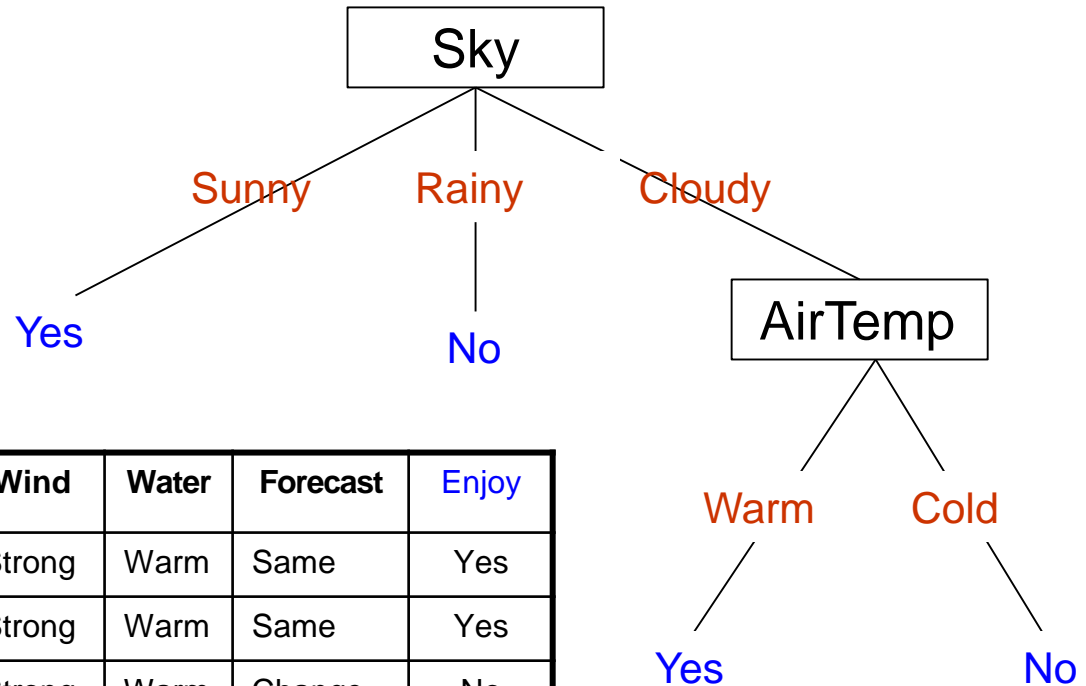
2-1 → 2.5 (Chapter 2, ML textbook)

Chapter 9 of the Vietnamese Textbook

Decision Trees


Example	Sky	AirTemp	Humidity	Wind	Water	Forecast	EnjoySport
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Sunny	Warm	High	Strong	Warm	Same	Yes
3	Rainy	Cold	High	Strong	Warm	Change	No
4	Sunny	Warm	High	Strong	Cool	Change	Yes
5	Cloudy	Warm	High	Weak	Cool	Same	Yes
6	Cloudy	Cold	High	Weak	Cool	Same	No

Decision Trees




No.	Sky	AirTemp	Humidity	Wind	Water	Forecast	Enjoy
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Sunny	Warm	High	Strong	Warm	Same	Yes
3	Rainy	Cold	High	Strong	Warm	Change	No
4	Sunny	Warm	High	Strong	Cool	Change	Yes
5	Cloudy	Warm	High	Weak	Cool	Same	Yes
6	Cloudy	Cold	High	Weak	Cool	Same	No

Decision Trees



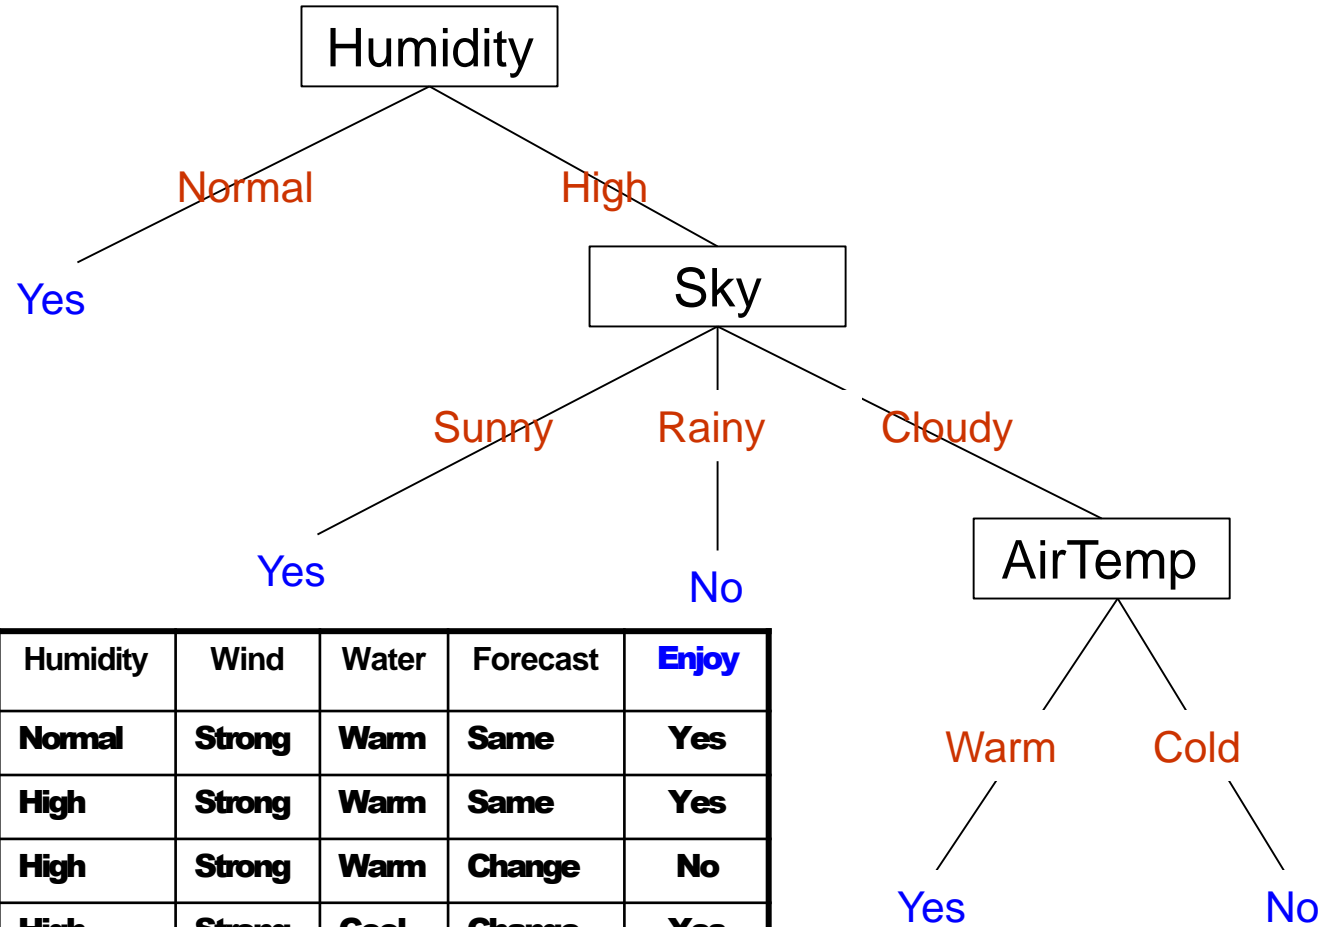
$(\text{Sky} = \text{Sunny}) \vee (\text{Sky} = \text{Cloudy} \wedge \text{AirTemp} = \text{Warm})$

Decision Trees



7	Rainy	Warm	Normal	Weak	Cool	Same	?
8	Cloudy	Warm	High	Strong	Cool	Change	?


Decision Trees




No.	Sky	AirTemp	Humidity	Wind	Water	Forecast	Enjoy
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Sunny	Warm	High	Strong	Warm	Same	Yes
3	Rainy	Cold	High	Strong	Warm	Change	No
4	Sunny	Warm	High	Strong	Cool	Change	Yes
5	Cloudy	Warm	High	Weak	Cool	Same	Yes
6	Cloudy	Cold	High	Weak	Cool	Same	No

Decision Trees

$$A_1 = v_1$$




$$A_2 = v_2$$



Homogeneity of Examples

- Entropy(S) = $-p_+ \log_2 p_+ - p_- \log_2 p_-$




Homogeneity of Examples

- Entropy(S) = $\sum_{i=1,c} -p_i \log_2 p_i$ impurity measure

Information Gain

- $\text{Gain}(S, A) = \text{Entropy}(S) - \sum_{v \in \text{Values}(A)} (|S_v|/|S|) \cdot \text{Entropy}(S_v)$




Example

- $\text{Entropy}(S) = -p_+ \log_2 p_+ - p_- \log_2 p_- = - (4/6) \log_2 (4/6) - (2/6) \log_2 (2/6)$
 $= 0.389 + 0.528 = 0.917$
- $\text{Gain}(S, \text{Sky})$
 $= \text{Entropy}(S) - \sum_{v \in \{\text{Sunny, Rainy, Cloudy}\}} (|S_v|/|S|) \text{Entropy}(S_v)$
 $= \text{Entropy}(S) - [(3/6) \cdot \text{Entropy}(S_{\text{Sunny}}) + (1/6) \cdot \text{Entropy}(S_{\text{Rainy}}) +$
 $(2/6) \cdot \text{Entropy}(S_{\text{Cloudy}})]$
 $= \text{Entropy}(S) - (2/6) \cdot \text{Entropy}(S_{\text{Cloudy}})$
 $= \text{Entropy}(S) - (2/6)[- (1/2) \log_2 (1/2) - (1/2) \log_2 (1/2)]$
 $= 0.917 - 0.333 = 0.584$

Example

- $\text{Entropy}(S) = -p_+ \log_2 p_+ - p_- \log_2 p_- = - (4/6) \log_2 (4/6) - (2/6) \log_2 (2/6)$
 $= 0.389 + 0.528 = 0.917$
- $\text{Gain}(S, \text{Water})$
 $= \text{Entropy}(S) - \sum_{v \in \{\text{Warm}, \text{Cool}\}} (|S_v|/|S|) \text{Entropy}(S_v)$
 $= \text{Entropy}(S) - [(3/6) \cdot \text{Entropy}(S_{\text{Warm}}) + (3/6) \cdot \text{Entropy}(S_{\text{Cool}})]$
 $= \text{Entropy}(S) - (3/6) \cdot 2 \cdot [- (2/3) \log_2 (2/3) - (1/3) \log_2 (1/3)]$
 $= \text{Entropy}(S) - 0.389 - 0.528$
 $= 0$

Example



- $\text{Gain}(S_{\text{Cloudy}}, \text{AirTemp})$
 $= \text{Entropy}(S_{\text{Cloudy}}) - \sum_{v \in \{\text{Warm, Cold}\}} (|S_v|/|S|) \text{Entropy}(S_v)$
 $= 1$
- $\text{Gain}(S_{\text{Cloudy}}, \text{Humidity})$
 $= \text{Entropy}(S_{\text{Cloudy}}) - \sum_{v \in \{\text{Normal, High}\}} (|S_v|/|S|) \text{Entropy}(S_v)$
 $= 0$

Inductive Bias

- Hypothesis space: complete!

Inductive Bias

- Hypothesis space: complete!
- Shorter trees are preferred over larger trees
- Prefer the simplest hypothesis that fits the data

Inductive Bias

- Decision Tree algorithm: searches **incompletely** thru a **complete** hypothesis space.

⇒ **Preference** bias

- Candidate-Elimination searches **completely** thru an **incomplete** hypothesis space.

⇒ **Restriction** bias


Overfitting

- $h \in H$ is said to **overfit** the training data if there exists $h' \in H$, such that h has smaller error than h' over the **training** examples, but h' has a smaller error than h over the **entire distribution** of instances.

Overfitting

- $h \in H$ is said to **overfit** the training data if there exists $h' \in H$, such that h has smaller error than h' over the **training** examples, but h' has a smaller error than h over the **entire distribution** of instances:
 - There is noise in the data
 - The number of training examples is too small to produce a representative sample of the target concept

Overfitting



Homework

Exercises 3-1→3.4 (Chapter 3, ML textbook)