

Heuristic search

Vuong Ba Thinh

Department of Computer Science - CSE Faculty - HCMUT

vbthinh@hcmut.edu.vn

July 6, 2015

Acknowledgement

Most of these slides were either created by Prof. Cao Hoang Tru at University of Technology or else are modifications of his slides.

Outline

- 1 Generate-and-Test
- 2 Hill Climbing
- 3 Simulated Annealing
- 4 Best-First Search
- 5 Genetic Algorithm

Algorithm

- 1 Generate a possible solution
- 2 Test to see if this is actually a solution
- 3 Quit if a solution has been found. Otherwise, return to step 1

Generate-and-Test

- Acceptable for simple problem
- Inefficient for problems with large space

Generate-and-Test

- **Exhaustive** generate-and-test.
- **Heuristic** generate-and-test: not consider paths that seem unlikely to lead to a solution.
- **Plan** generate-and-test:
 - Create a list of candidates.
 - Apply generate-and-test to that list.

Example (colored blocks)

Arrange four 6-sided cubes in a row, with each side of each cube painted one of four colors, such that on all four sides of the row one block face of each color is showing.

Example (heuristic)

if there are more **red** faces than other colors then, when placing a block with several **red** faces, use few of them as possible as outside faces.

Hill Climbing

Searching for a **goal state** = Climbing to the **top of a hill**

Hill Climbing

- Generate-and-Test + **direction to move**.
- **Heuristic function**: estimate how close a given state to a goal state

Algorithm

- ① Evaluate the initial state
- ② Loop until a solution is found or there are no new operators left to be applied:
 - Select and apply a operator
 - Evaluate the new state:
 - goal \rightarrow quit
 - better than current state \rightarrow new current state

Algorithm

- ① Evaluate the initial state
- ② Loop until a solution is found or there are no new operators left to be applied:
 - Select and apply a operator
 - Evaluate the new state:
 - goal \rightarrow quit
 - better than current state \rightarrow new current state

not try all possible new states!

Simple Hill Climbing

Example (colored blocks)

Heuristic function: the sum of the number of different colors on each of the four sides (solution = 16).

Heuristic function

Heuristic function as a way to inject **task-specific knowledge** into the control process.

Steepest-Ascent Hill Climbing (Gradient Search)

- Considers **all the moves** from the current state.
- Selects **the best one** as the next state.

Steepest-Ascent Hill Climbing (Gradient Search)


Algorithm

- ① Evaluate the initial state
- ② Loop until a solution is found or there are no new operators left to be applied:
 - Apply all the possible operators
 - Evaluate **the best new state**:
 - goal \rightarrow quit
 - better than current state \rightarrow new current state

Hill Climbing: Disadvantages

Local maximum

A state that is better than all of its neighbors, but not better than some other states far away.



Hill Climbing: Disadvantages

Plateau


A flat area of the search space in which all neighboring states have the same value.



Hill Climbing: Disadvantages

Ridge

The orientation of the high region, compared to the set of available moves, makes it impossible to climb up. However, two moves executed serially may increase the height.



Hill Climbing: Disadvantages


Ways Out

- **Backtrack** to some earlier node and try going in a different direction.
- Make a **big jump** to try to get in a new section.
- Moving in **several directions** at once.

Hill Climbing: Disadvantages

- Hill climbing is a **local method**: Decides what to do next by looking only at the “immediate” consequences of its choices.
- **Global information** might be encoded in heuristic functions.

Hill Climbing: Block World




Hill Climbing: Block World




Local heuristic:

- **+1** for each block that is resting on the thing it is supposed to be resting on.
- **-1** for each block that is resting on a wrong thing.

Hill Climbing: Block World



Hill Climbing: Block World




Hill Climbing: Block World



Global heuristic:

- For each block that has the correct support structure: **+1** to every block in the support structure.
- For each block that has a wrong support structure: **-1** to every block in the support structure.

Hill Climbing: Block World



Hill Climbing: Conclusion

- Can be very **inefficient** in a large, rough problem space.
- Global heuristic may have to **pay for computational complexity**.
- **Often useful** when combined with other methods, getting it started right in the right general neighborhood.

Simulated Annealing

A variation of hill climbing in which, at the beginning of the process, some downhill moves may be made.

Simulated Annealing

- To do **enough exploration of the whole space** early on, so that the final solution is relatively insensitive to the starting state.
- **Lowering the chances** of getting caught at a local maximum, or plateau, or a ridge.

Physical Annealing


- Physical substances are melted and then **gradually cooled** until some solid state is reached.
- The goal is to produce a **minimal-energy state**.
- **Annealing schedule**: if the temperature is lowered sufficiently slowly, then the goal will be attained.
- Nevertheless, there is some **probability** for a transition to a higher energy state: $e^{-\Delta E/kT}$

Algorithm

- ① Evaluate the initial state
- ② Loop until a solution is found or there are no new operators left to be applied:
 - Set T according to an annealing schedule.
 - Selects and applies a new operator
 - Evaluate the new state:
 - goal \rightarrow quit
 - $\Delta E = \text{Val}(\text{current state}) - \text{Val}(\text{new state})$
 - $\Delta E < 0 \rightarrow$ new current state
 - else \rightarrow new current state with probability $e^{-\Delta E/kT}$

Depth-first search


- Pros: not having to expand all competing branches
- Cons: getting trapped on dead-end paths



Best-First Search


Breadth-first search

- Pros: not getting trapped on dead-end paths
- Cons: having to expand all competing branches







Best-First Search

⇒ Combining the two is to **follow a single path at a time**, but **switch paths** whenever some competing path looks more promising than the current one.



Best-First Search



Best-First Search

- **OPEN**: nodes that have been generated, but have not examined.
This is organized as a **priority queue**.
- **CLOSED**: nodes that have already been examined.
Whenever a new node is generated, **check** whether it has been **generated before**.

Algorithm

- ① $OPEN = \{\text{initial state}\}.$
- ② Loop until a goal is found or there are no nodes left in OPEN:
 - Pick the best node in OPEN
 - Generate its successors
 - For each successor:
 - new \rightarrow evaluate it, add it to OPEN, record its parent generated before
 - \rightarrow change parent, update successors

Best-First Search

- Greedy search:

$h(n)$ = cost of the cheapest path from node n to a goal state.

Best-First Search

- Greedy search:

$h(n)$ = cost of the cheapest path from node n to a goal state.

- Uniform-cost search:

$g(n)$ = cost of the cheapest path from the initial state to node n .

Best-First Search

- Greedy search:

$h(n)$ = cost of the cheapest path from node n to a goal state.

Neither optimal nor complete

- Uniform-cost search:

$g(n)$ = cost of the cheapest path from the initial state to node n .

Best-First Search

- Greedy search:

$h(n)$ = cost of the cheapest path from node n to a goal state.

Neither optimal nor complete

- Uniform-cost search:

$g(n)$ = cost of the cheapest path from the initial state to node n .

Optimal and complete, but very inefficient

Best-First Search

Algorithm: A^* (Hart et al., 1968)

$$f(n) = g(n) + h(n)$$

$h(n)$ = cost of the cheapest path from node n to a **goal state**.

$g(n)$ = cost of the cheapest path from the **initial state** to node n .

Best-First Search

Algorithm: A^* (Hart et al., 1968)

$$f^*(n) = g^*(n) + h^*(n)$$

$h^*(n)$ (heuristic factor) = estimate of $h(n)$.

$g^*(n)$ (depth factor) = approximation of $g(n)$ found by A^* so far.

Genetic Algorithm



Algorithm: Gradient search



Algorithm: Genetic algorithm

- A **genetic algorithm** is a heuristic search that **mimics** the process of **natural evolution**.
- There are **five phases**:
 - Initial Population
 - Fitness function
 - Selection
 - Crossover
 - Mutation

Genetic Algorithm: Initial Population

- Identify the information already presents in the problem as well as information that needs to be computed or derived
- Encode the solution in a form that is analogous to biological's chromosomes or sequence of DNA
- Generate randomly states that satisfy the problem

Genetic Algorithm: Initial Population

Example (8-queens)

The eight queens puzzle is the problem of placing eight chess queens on an 8x8 chessboard so that no two queens threaten each other

Representation

- Each state must have 8 queens
- One queen in each column
- Usually represented by a bit-string

EX: [1,2,3,4,5,6,7,8]

Genetic Algorithm: Initial Population




Figure: [4,2,7,3,6,8,1,5]

24748552

32752411


24415124

32543213

Initial Population

Genetic Algorithm: Fitness function

- Evaluate the “goodness” of each candidate solution
- The probability of being chosen for reproduction is based on your fitness score
- EX: 8-queens → fitness function: the number of nonattacking pairs of queens



Genetic Algorithm: Selection

- Fitness Proportionate Selection

based on probability of selection: $p_i = f_i / \sum f_i$

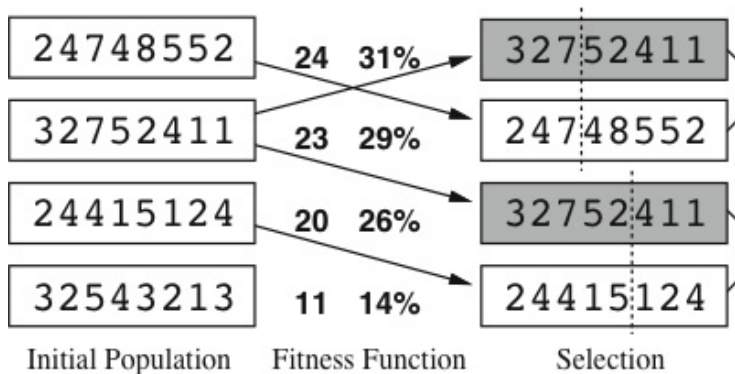
- Tournament Selection

based on pre-determine number T (tournament size)

- Rank Selection

rank the individuals in the population in ascending order according to their fitness values. Then, select similar to “Fitness Proportionate Selection”

Genetic Algorithm: Selection




Genetic Algorithm: Crossover

Single-point Crossover


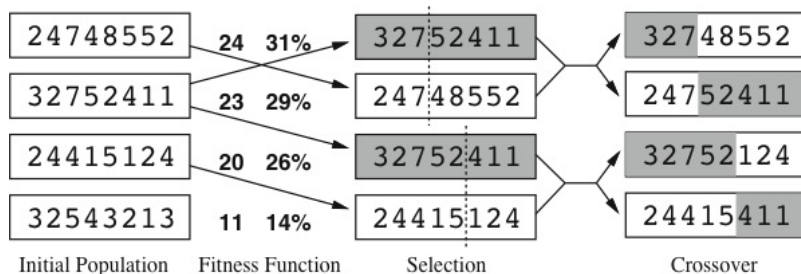


Multi-point Crossover




- Not all chosen chromosome pairs will undergo crossover
- p_c : probability of crossover or crossover rate

Genetic Algorithm: Crossover




Genetic Algorithm: Mutation

- Mutation **replaces the values** of some randomly chosen genes of a chromosome by some arbitrary new values
- **Bit inversion**
- p_m : **probability of mutation** or **mutation rate**



Genetic Algorithm: Mutation



- 8-puzzle
- Graph-coloring
- 8-queens
- Sudoku